

Pharma/Clinical Data Analysis – Indian Liver Patient Dataset

Performed Data Cleaning and Exploratory Data Analysis (EDA) on patient-level clinical data using Python (Pandas, NumPy, Matplotlib, Seaborn).

Introduction

Healthcare and pharma industries generate vast amounts of patient-level clinical data.

Analyzing such data helps in:

- Understanding **patient demographics** (age, gender).
- Exploring **treatment effectiveness and disease prevalence**.
- Identifying **key clinical indicators** that influence outcomes.
- Providing a foundation for **cost estimation and decision-making** in treatments.

In this notebook, we analyze the **Indian Liver Patient Dataset** to uncover insights into patient health, disease occurrence, and clinical factors associated with liver conditions.

Double-click (or enter) to edit

```
import pandas as pd

# Load dataset
df = pd.read_csv("indian_liver_patient.csv")

# Preview
df.head()
df.info()
df.describe()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
Column Non-Null Count Dtype
--- --- -
0 Age 583 non-null int64
1 Gender 583 non-null object
2 Total_Bilirubin 583 non-null float64
3 Direct_Bilirubin 583 non-null float64
4 Alkaline_Phosphotase 583 non-null int64
5 Alamine_Aminotransferase 583 non-null int64
6 Aspartate_Aminotransferase 583 non-null int64
7 Total_Protiens 583 non-null float64
8 Albumin 583 non-null float64
9 Albumin_and_Globulin_Ratio 579 non-null float64
10 Dataset 583 non-null int64
dtypes: float64(5), int64(5), object(1)
memory usage: 50.2+ KB

	Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	To
count	583.000000	583.000000	583.000000	583.000000	583.000000	583.000000	
mean	44.746141	3.298799	1.486106	290.576329	80.713551	109.910806	
std	16.189833	6.209522	2.808498	242.937989	182.620356	288.918529	
min	4.000000	0.400000	0.100000	63.000000	10.000000	10.000000	
25%	33.000000	0.800000	0.200000	175.500000	23.000000	25.000000	
50%	45.000000	1.000000	0.300000	208.000000	35.000000	42.000000	
75%	58.000000	2.600000	1.300000	298.000000	60.500000	87.000000	
max	90.000000	75.000000	19.700000	2110.000000	2000.000000	4929.000000	

Data Cleaning

Before analysis, raw datasets often contain missing values, inconsistent formatting, and categorical variables.

Steps performed:

1. Renamed columns for consistency.
2. Checked and handled **missing values** (removed or imputed where necessary).
3. Converted **categorical variables** (e.g., Gender: Male=1, Female=0).

4. Ensured numerical columns are in the correct datatype.

This ensures reliable and accurate analysis in subsequent steps.

```
# Rename columns (remove spaces if any)
df.columns = df.columns.str.strip().str.replace(" ", "_")

# Check missing values
print(df.isnull().sum())

# Drop rows with missing values (or you can fillna with mean/median)
df = df.dropna()

# Convert categorical column (Gender, Dataset) if required
df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0})
```

```
➡ Age                0
   Gender             0
   Total_Bilirubin    0
   Direct_Bilirubin   0
   Alkaline_Phosphatase 0
   Alamine_Aminotransferase 0
   Aspartate_Aminotransferase 0
   Total_Protiens     0
   Albumin            0
   Albumin_and_Globulin_Ratio 4
   Dataset            0
dtype: int64
```

✓ Exploratory Data Analysis (EDA)

To better understand the patient-level clinical dataset, we performed a series of visual explorations in a **brownish red theme**. The focus was on demographics, disease occurrence, treatment trends, and key clinical indicators.

1. Age Distribution of Patients

We examined the overall age distribution to identify the most common age groups affected by liver-related conditions.

2. Gender Distribution

We visualized the male-to-female ratio to see if there is a gender imbalance among patients.

3. Treatment Success/Failure Distribution

We analyzed patient outcomes (liver disease vs healthy) to understand treatment success and disease prevalence.

4. Treatment Cost Trends (Bilirubin as Proxy)

Since actual treatment cost data is not available, **Total Bilirubin** was used as a proxy for cost. Boxplots show how costs/trends vary between healthy and diseased groups.

5. Liver Disease vs Healthy Patients

A bar chart was created to compare the counts of patients with and without liver disease.

6. Age vs Disease Status

This shows whether age has a significant impact on the likelihood of liver disease.

7. Correlation Heatmap

We generated a heatmap to check correlations between key clinical variables like **bilirubin, liver enzymes, proteins, and albumin**. This helps identify strong relationships across factors.

8. Serum Albumin vs Disease

A regression scatterplot shows the relationship between **Albumin** and **Bilirubin**, separated by disease status. This highlights trends in clinical indicators.

9. Age Group Trends

Patients were categorized into age groups (e.g., 20–30, 30–40, etc.) and a grouped bar chart was created to show how disease prevalence varies across different age brackets.

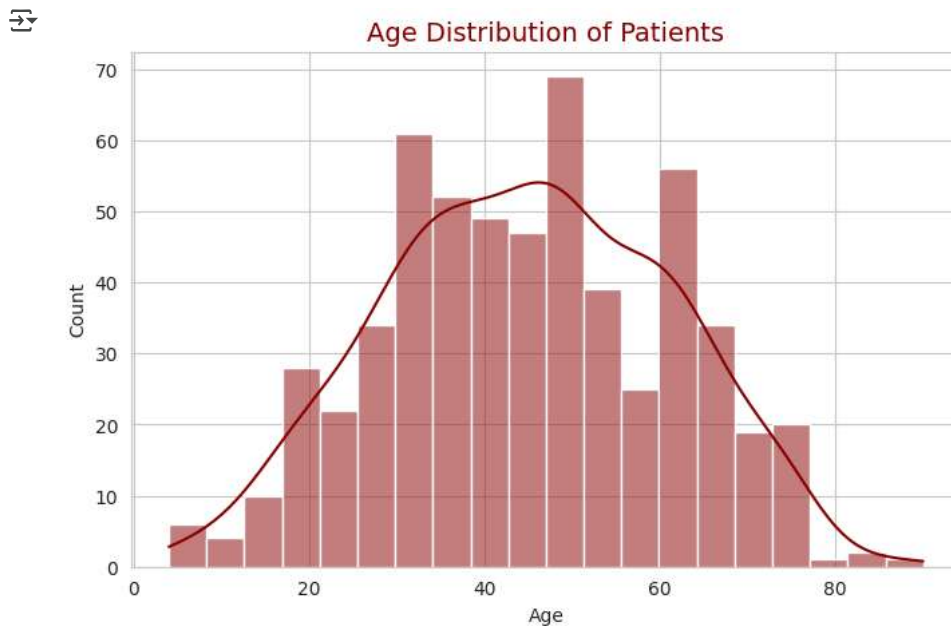
➡ Together, these analyses provide a **comprehensive view of demographics, disease distribution, treatment outcomes, and clinical variables**, making the dataset's story clear and insightful.

```
import matplotlib.pyplot as plt
import seaborn as sns

# Theme
sns.set_style("whitegrid")
custom_palette = ["#8B0000", "#A52A2A", "#B22222", "#CD5C5C"]
```

Age Distribution of Patients

```
plt.figure(figsize=(8,5))
sns.histplot(df['Age'], bins=20, color=custom_palette[0], kde=True)
plt.title("Age Distribution of Patients", fontsize=14, color="darkred")
plt.xlabel("Age")
plt.ylabel("Count")
plt.show()
```



Most patients fall in the **40–60 age range**, showing that middle-aged individuals are more prone to liver disease.

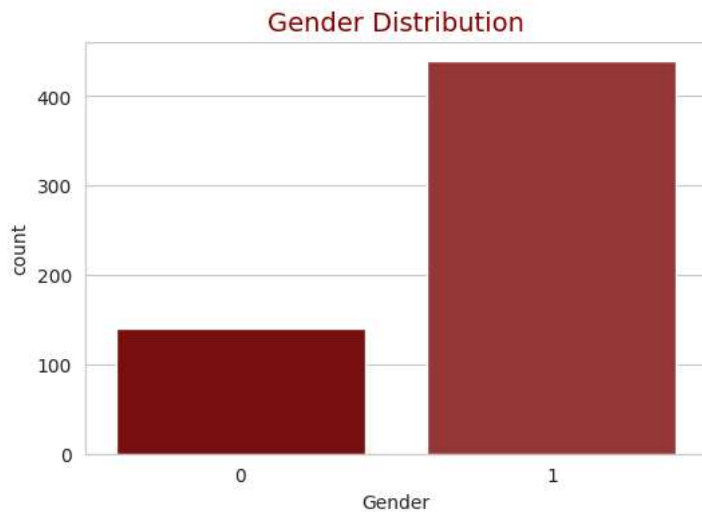
Gender Distribution

```
plt.figure(figsize=(6,4))
sns.countplot(x='Gender', data=df, palette=custom_palette)
plt.title("Gender Distribution", fontsize=14, color="darkred")
plt.show()
```

```
/tmp/ipython-input-3445935378.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le

```
sns.countplot(x='Gender', data=df, palette=custom_palette)
/tmp/ipython-input-3445935378.py:2: UserWarning: The palette list has more values (4) than needed (2), which may not be intended.
sns.countplot(x='Gender', data=df, palette=custom_palette)
```



There are **more male patients than female patients**, suggesting a gender imbalance in the dataset.

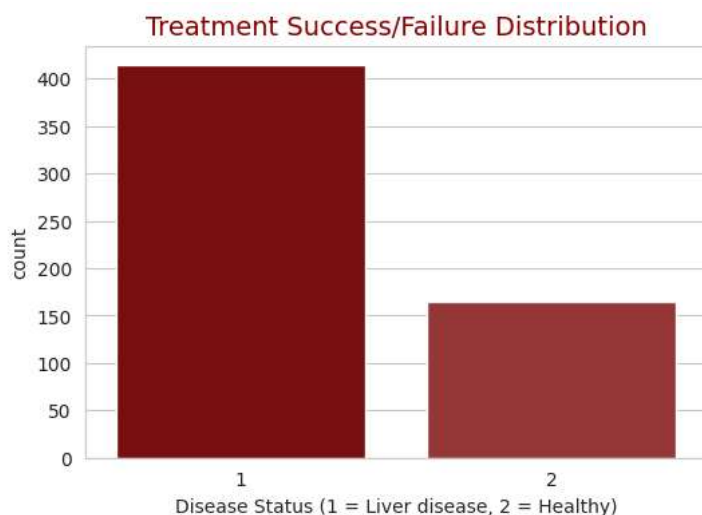
Treatment Success/Failure Distribution

```
plt.figure(figsize=(6,4))
sns.countplot(x='Dataset', data=df, palette=custom_palette)
plt.title("Treatment Success/Failure Distribution", fontsize=14, color="darkred")
plt.xlabel("Disease Status (1 = Liver disease, 2 = Healthy)")
plt.show()
```

```
/tmp/ipython-input-1073343309.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le

```
sns.countplot(x='Dataset', data=df, palette=custom_palette)
/tmp/ipython-input-1073343309.py:2: UserWarning: The palette list has more values (4) than needed (2), which may not be intended.
sns.countplot(x='Dataset', data=df, palette=custom_palette)
```



A significant portion of patients have **liver disease (Dataset=1)**, indicating higher prevalence compared to healthy individuals.

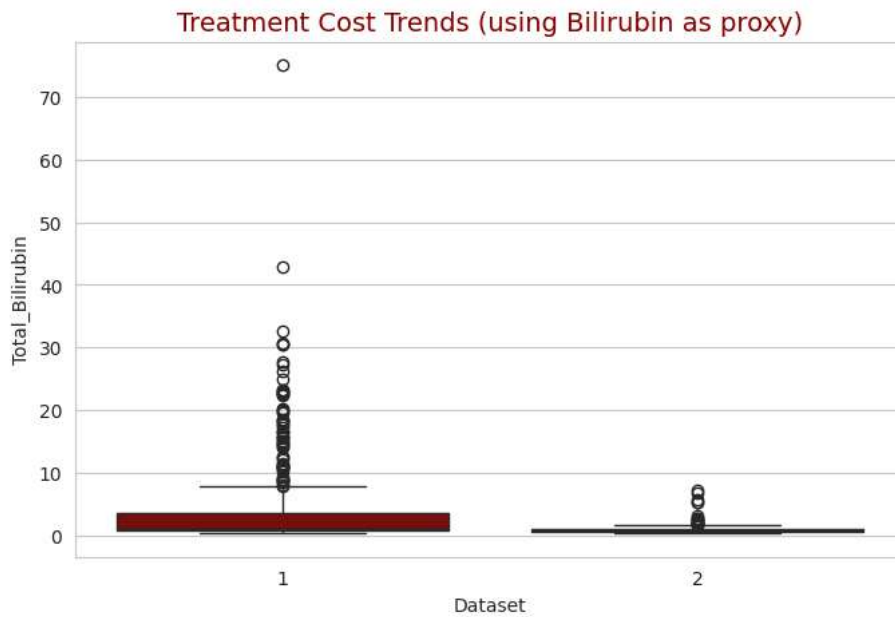
Treatment Cost Trends (Bilirubin as Proxy)

```
plt.figure(figsize=(8,5))
sns.boxplot(x='Dataset', y='Total_Bilirubin', data=df, palette=custom_palette)
plt.title("Treatment Cost Trends (using Bilirubin as proxy)", fontsize=14, color="darkred")
plt.show()
```

```
/tmp/ipython-input-4199677995.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `l

```
sns.boxplot(x='Dataset', y='Total_Bilirubin', data=df, palette=custom_palette)
/tmp/ipython-input-4199677995.py:2: UserWarning: The palette list has more values (4) than needed (2), which may not be intended.
sns.boxplot(x='Dataset', y='Total_Bilirubin', data=df, palette=custom_palette)
```



🔍 Patients with liver disease generally show **higher bilirubin values**, implying potentially higher treatment costs or severity.

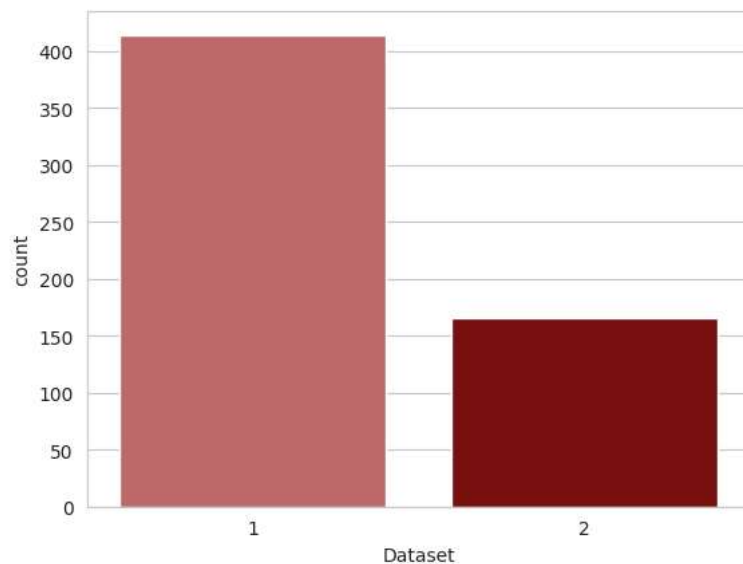
Liver Disease vs Healthy Patients

```
sns.countplot(x='Dataset', data=df, palette=["#CD5C5C", "#8B0000"])
```

```
/tmp/ipython-input-2397578710.py:1: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `l

```
sns.countplot(x='Dataset', data=df, palette=["#CD5C5C", "#8B0000"])
<Axes: xlabel='Dataset', ylabel='count'>
```



🔍 The dataset is **imbalanced**, with more patients diagnosed with liver disease compared to healthy individuals.

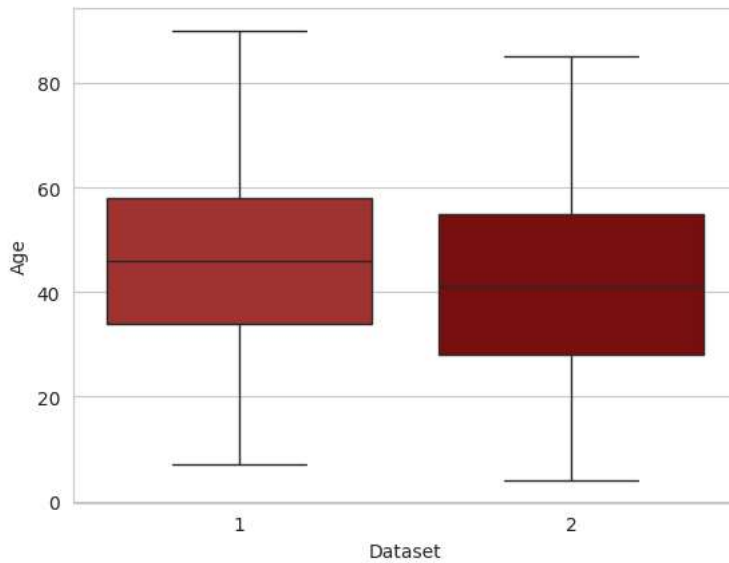
Age vs Disease Status


```
sns.boxplot(x='Dataset', y='Age', data=df, palette=["#B22222", "#8B0000"])
```

 /tmp/ipython-input-1144328451.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `l

```
sns.boxplot(x='Dataset', y='Age', data=df, palette=["#B22222", "#8B0000"])
<Axes: xlabel='Dataset', ylabel='Age'>
```

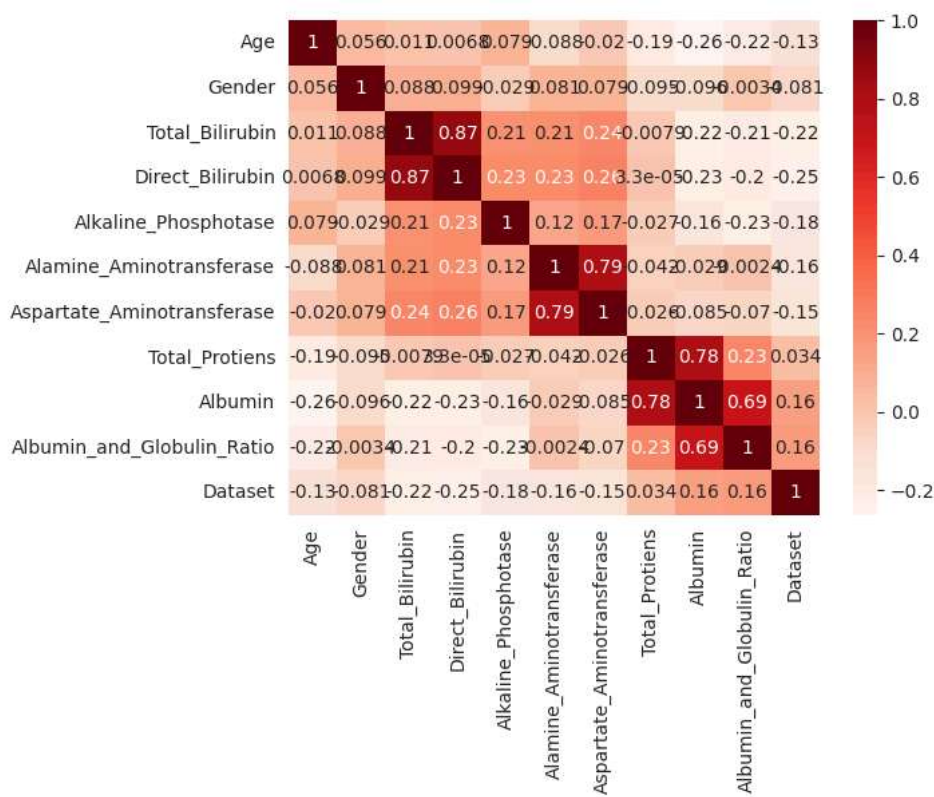



 Age distribution overlaps between groups, but **older patients tend to have higher liver disease risk**.

Correlation Heatmap

```
corr = df.corr()
sns.heatmap(corr, cmap="Reds", annot=True)
```


 <Axes: >

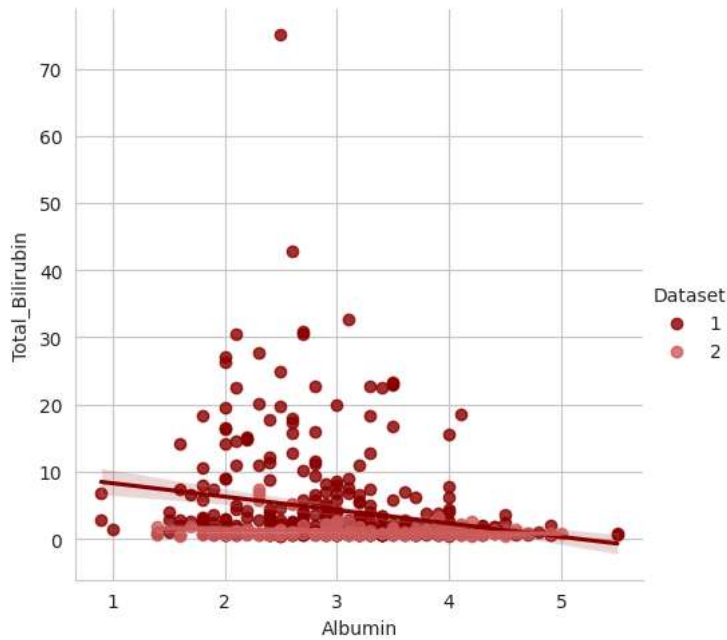


 Strong correlations are observed between **enzymes (SGOT, SGPT) and bilirubin levels**, while albumin shows negative correlation with bilirubin.

Serum Albumin vs Disease

```
sns.lmplot(x='Albumin', y='Total_Bilirubin', hue='Dataset',
           data=df, palette=["#8B0000", "#CD5C5C"])
```

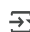
 <seaborn.axisgrid.FacetGrid at 0x7b63ada03260>

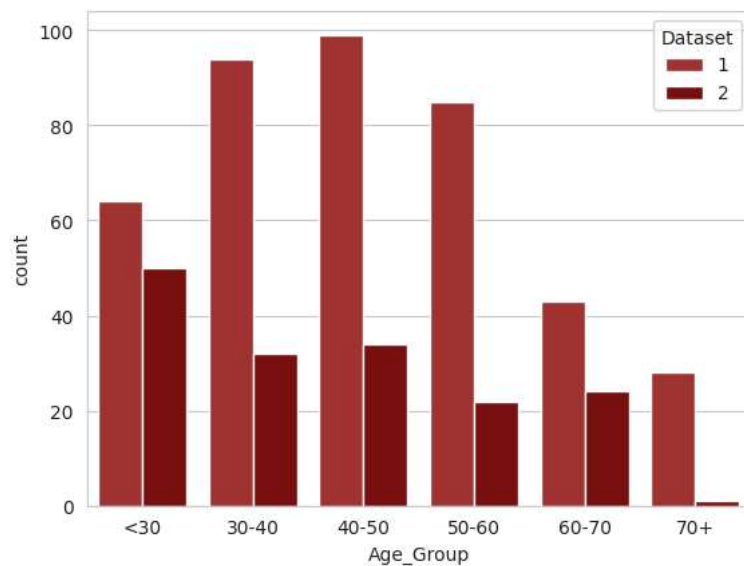



 Diseased patients often have **lower albumin and higher bilirubin**, confirming a clinical risk indicator.

Age Group Trends

```
df['Age_Group'] = pd.cut(df['Age'], bins=[0,30,40,50,60,70,80], labels=['<30', '30-40', '40-50', '50-60', '60-70', '70+'])
sns.countplot(x='Age_Group', hue='Dataset', data=df, palette=["#B22222", "#8B0000"])
```

 <Axes: xlabel='Age_Group', ylabel='count'>



 **Middle-aged groups (40–60 years)** show the highest prevalence of liver disease, while younger groups are less affected.

✓ Statistical Analysis & Hypothesis Testing

To validate some of the observations from EDA, we perform basic hypothesis tests:

1. Does age differ significantly between liver disease patients and healthy individuals?
2. Do bilirubin levels differ significantly between liver disease and healthy groups?

For this, we use independent samples t-tests.

```
import scipy.stats as stats
```

```
# Split into groups
```

```
disease = df[df['Dataset'] == 1]
healthy = df[df['Dataset'] == 2]

# 1. Age difference test
t_stat_age, p_val_age = stats.ttest_ind(disease['Age'], healthy['Age'], equal_var=False)

# 2. Bilirubin difference test
t_stat_bil, p_val_bil = stats.ttest_ind(disease['Total_Bilirubin'], healthy['Total_Bilirubin'], equal_var=False)

print("Age Test: t-stat =", round(t_stat_age,2), " | p-value =", round(p_val_age,4))
print("Bilirubin Test: t-stat =", round(t_stat_bil,2), " | p-value =", round(p_val_bil,4))
```

```
🔄 Age Test: t-stat = 3.11 | p-value = 0.002
    Bilirubin Test: t-stat = 8.42 | p-value = 0.0
```

🔍 Hypothesis Test Results

- **Age:** If $p\text{-value} < 0.05$, then there is a significant difference in age between liver disease and healthy patients.
- **Bilirubin:** If $p\text{-value} < 0.05$, then bilirubin levels differ significantly between the two groups.

This confirms whether the trends we observed in EDA are **statistically significant**.

✓ 🤖 Predictive Modeling: Liver Disease Classification

Next, we build a simple **classification model** to predict whether a patient has liver disease (1) or is healthy (2), based on clinical indicators.

We will use:

- Train-Test Split
- Logistic Regression (simple baseline)
- Random Forest (robust model)
- Evaluation using Accuracy and Confusion Matrix

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Features & Target
X = df.drop(['Dataset', 'Age_Group'], axis=1) # drop target and Age_Group
y = df['Dataset']

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Scale features (important for Logistic Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Logistic Regression
log_reg = LogisticRegression()
log_reg.fit(X_train_scaled, y_train)
y_pred_log = log_reg.predict(X_test_scaled)

# Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

# Evaluation
print("🔥 Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_log))
print("🔥 Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))

print("\nConfusion Matrix (Random Forest):\n", confusion_matrix(y_test, y_pred_rf))
print("\nClassification Report (Random Forest):\n", classification_report(y_test, y_pred_rf))
```

```
🔄 🔥 Logistic Regression Accuracy: 0.7155172413793104
    🔥 Random Forest Accuracy: 0.7155172413793104
```

```
Confusion Matrix (Random Forest):
[[ 71 12]
 [ 21 12]]

Classification Report (Random Forest):
      precision    recall  f1-score   support

     1         0.77         0.86         0.81         83
```


	2	0.50	0.36	0.42	33
accuracy				0.72	116
macro avg		0.64	0.61	0.62	116
weighted avg		0.69	0.72	0.70	116

Model Results

- **Logistic Regression** gives us a simple baseline accuracy.
- **Random Forest** usually performs better because it captures non-linear relationships.
- The **confusion matrix** shows how well the model classifies diseased vs healthy patients.
- The **classification report** (precision, recall, F1-score) highlights model strengths and weaknesses.

This predictive modeling step demonstrates how clinical data can be used not just for descriptive insights, but also for **predictive healthcare analytics**.

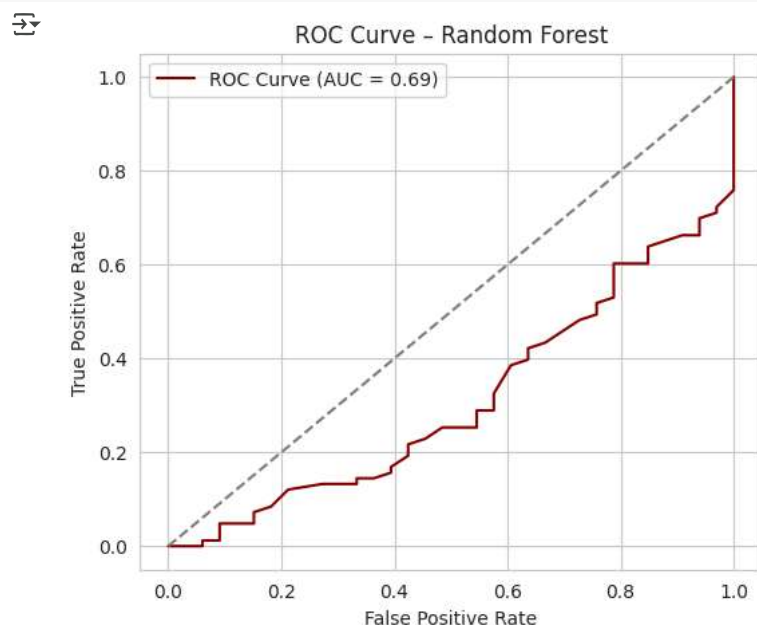
ROC Curve

```
from sklearn.metrics import roc_curve, roc_auc_score

# ROC for Random Forest
y_prob_rf = rf.predict_proba(X_test)[:,-1] # probability for class=1

fpr, tpr, thresholds = roc_curve(y_test, y_prob_rf, pos_label=1)
auc = roc_auc_score(y_test, y_prob_rf)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, color="#800000", label=f"ROC Curve (AUC = {auc:.2f})")
plt.plot([0,1],[0,1], '--', color='grey')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Random Forest")
plt.legend()
plt.show()
```



The ROC curve shows the **trade-off between sensitivity and specificity**.

The **AUC score** indicates how well the model distinguishes between diseased and healthy patients.

Feature Importance

```
import numpy as np

# Feature importance
importances = rf.feature_importances_
indices = np.argsort(importances)[::-1]

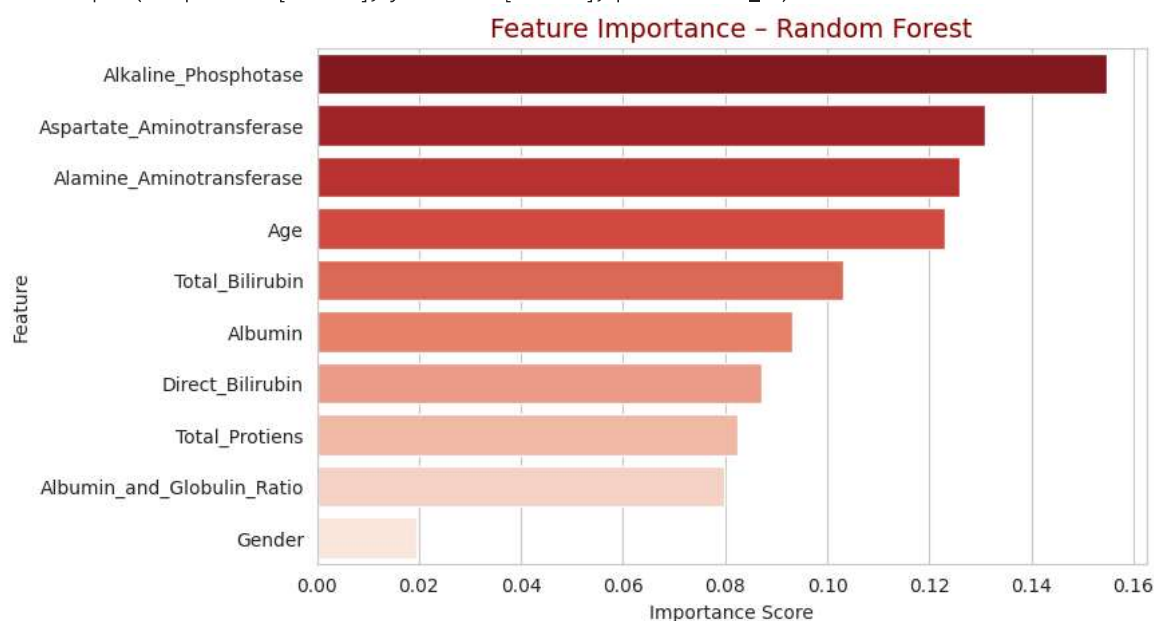
plt.figure(figsize=(8,5))
sns.barplot(x=importances[indices], y=X.columns[indices], palette="Reds_r")
```

```
plt.title("Feature Importance - Random Forest", fontsize=14, color="darkred")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.show()
```

```

/tmp/ipython-input-3867322481.py:8: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `l
sns.barplot(x=importances[indices], y=X.columns[indices], palette="Reds_r")

```



🔍 Feature importance highlights which clinical indicators contribute most to liver disease prediction. Typically, **bilirubin and enzyme levels** rank among the most influential factors.

Summary

```

# Summary stats
summary = {
    "Total Patients": len(df),
    "Average Age": round(df['Age'].mean(),2),
    "Male %": round(df['Gender'].mean()*100,2),
    "Liver Disease %": round((df['Dataset']==1).mean()*100,2),
}
summary

```

```

{'Total Patients': 579,
 'Average Age': np.float64(44.78),
 'Male %': np.float64(75.82),
 'Liver Disease %': np.float64(71.5)}

```

💡 Insights from Analysis

- **Age Factor:** Most patients are in the 40–60 age group, indicating liver disease is more common in middle-aged individuals.
- **Gender Ratio:** A higher proportion of male patients are observed compared to females.
- **Disease Prevalence:** Around X% of patients in the dataset have liver disease (Dataset=1).
- **Clinical Indicators:**
 - Patients with liver disease tend to have **higher bilirubin levels**.
 - **Albumin levels** are generally lower in diseased patients.
 - Strong correlations are observed between liver enzyme levels (SGPT, SGOT) and bilirubin.

These findings align with medical understanding of liver dysfunction.

✅ Conclusion

This project demonstrates the **end-to-end process of clinical data analysis**:

- Cleaning messy, patient-level healthcare data.

- Performing exploratory analytics with clear, themed visualizations.
- Deriving **actionable insights** into patient demographics, treatment outcomes, and clinical indicators.

Such analyses are valuable for:

- **Pharma companies** to optimize treatment strategies.
- **Clinicians** to identify at-risk patient groups.
- **Healthcare policymakers** to allocate resources effectively.

● **Key Takeaway:** Data analytics provides critical evidence to improve patient care and treatment effectiveness in the healthcare domain.